

Computer Security - CheatSheet

IN BA5 - Thomas Bourgeat

Notes by Ali EL AZDI

CompSec Properties - Confidentiality. prevent unauthorized disclosure of information. <i>authorized users may read a file</i>	Harm. The bad thing that could happen when the threat materializes. (<i>adversary steals the money, learns my password...</i>)
- Integrity. prevent unauthorized modification of information. <i>authorized programs may write a file</i>	Security Policy. high level description of the security properties that must hold in the system in relation to assets and principals.
- Availability. prevent unauthorized denial of service or access to information and resources. <i>authorized services can access a file</i>	- Assets (objects). anything with value (data, files, memory) needing protection.
- Authenticity. assurance that entities (users, systems, or data) are genuine and can be verified as such.	- Principals (subjects). people, computer programs, services.
- Anonymity. protection of an individual's identity from being disclosed or linked to specific actions or data.	Security Mechanism. Technical mechanism used to ensure that the security policy is not violated by an adversary within the threat model, we can only prepare for threats we're aware of
- Non-repudiation. assurance that a party in a communication cannot deny the authenticity of their signature or the sending of a message.	(<i>Policy. ensure messages cannot be read by anyone but the sender and the receiver, Mechanism. encrypt the message before sending</i>)
The Adversary. malicious entity aiming at breaching the security policy and will choose the optimal way to use its resources to mount an attack that violates the security properties.	Composition of Security Mechanisms
Threat Model. describes the resources available to the adversary and their capabilities (<i>has access to internet, but doesn't have access to the internal network of the company.</i>)	- Defence in depth. As long as one remains unbroken the Security Policy isn't broken (<i>two-factor auth</i>)
Threat. Who might attack which assets, using what resources, with what goal, how, and with what probability	- Weakest Link. if anyone fails the Security Policy, it is broken. (<i>security questions for a lost password → just need to know the answer...</i>)
Vulnerability. Specific weakness that adversaries could exploit with interest in a lot of assets (<i>API is not protected, password appears in plain text...</i>)	<i>Humans can be vulnerabilities - phishing attacks, bad use of passwords...</i>

Principles of CompSec.

1. **Economy of mechanism** Keep security mechanism design simple and small ⇒ Easier to audit and verify, testing is **not** appropriate to evaluate security.

- **Trusted Computing Base (TCB).** Every component of the system on which the security policy relies upon *hardware, firmware, software*.

The TCB is trusted to operate correctly for the security policy to hold. → If something goes wrong in it, the security policy may be violated

must be kept small to easy verification / diminish the attack surface

2. **Fail-safe defaults.** Base access decisions on permission rather than exclusion. (*Whitelist, do not blacklist*)

If something fails, be as secure as it does not fail errors / uncertainty should error on the side of the security policy.

3. **Complete mediation..** Every access to every object must be checked for authority. A **Reference Monitor** mediates all actions from subjects on objects and ensures they are according to the policy. Δ time to check vs. time to use

4. **Open design** The design should not be secret *Always design as if the enemy knows the system. (Crypto-only secret is key. Authentication-Only secret is password, for Obfuscation-Only used secret is noise.)*. **assuming the threat model can't get a hold of the system is unrealistic (employee corruption, ...)**

Access Control. Security mechanism that ensures all accesses and actions on objects by principals are within the security policy.

no chicken soup (checks everywhere in code), use a reference monitor, module used all over code checking for subjects and actions

Discretionary Access Control (DAC). Object owners assign permissions, ownership of resources. *Linux, Social Networks*

Mandatory Access Control (MAC). Central security policy assigns permissions, usually for organizations with need for central control. *Military, Hospital,....*

Access Control Matrix. abstract representation of all permitted triplets of (subject, object, access right) within a system.

	file1	file2	file3
Alice	read, write		read
Bob		read, write	read, write

Complexity. $O(f \cdot u)$

Access Control List (ACL). associate permissions to objects, stores permissions close to the resource.

file1: (Alice,read/write), file2: (Bob, read/write),

file3: (Alice,read), (Bob,read/write)

⊕ easy to determine who can access a resource and to revoke rights by resource.

⊖ hard to check all users rights, to remove all perms. for a user, delegate perms.

Role Based Access Control (RBAC). access granted based on user roles and predefined permissions. systems have too many subjects (that come and go) → large dynamic ACLs. Subjects are often similar to each other and get assigned the same rights.

1. Assign permissions to roles, 2. Assign roles to subjects

3. Subjects have the permissions of their assigned role

⊖ role explosion (temptation to create fine grained roles), limited expressiveness, difficult to implement separation of privilege + least privilege.

Group Based Access Control (GBAC). access granted based on user group membership. **exactly like RBAC but instead of roles, permissions are assigned to groups.** groups are broader, represent organizational units instead than specific functions.

1. Assign permissions to groups 2. Assign subjects to groups

3. Subjects have the combined permissions of all their groups

⊖ coarse granularity, overlapping group memberships, inconsistent permissions, difficult to manage if users belong to many groups.

In case of **Negative Permissions** check negative permissions first before group, Δ system crashes before negative checks.

Ambient Authority. an action succeeds if the subject only specifies the *operation* and the *object name*, not the specific authority used.

⊖ leads to accidental misuse of authority, programs may act with more rights than intended. → Confused Deputy Problem.

Confused Deputy Problem. when a program (deputy) is tricked into misusing its authority on behalf of another subject.

Alice runs compiler(input, bill) ⇒ compiler (with write access to bill) overwrites billing file. ⇒ Alice uses compiler's authority to modify bill indirectly.

⊖ ambient authority allows unintended privilege use.

Solutions:

1. Restrict privileged process access.

2. Make privileged process check user's authorization.

3. Use **Capabilities** to explicitly delegate rights.

Harm. The bad thing that could happen when the **threat** materializes. (*adversary steals the money, learns my password...*)

Security Policy. high level description of the security properties that must hold in the system in relation to assets and principals.

- **Assets (objects).** anything with value (data, files, memory) needing protection.

- **Principals (subjects).** people, computer programs, services.

Security Mechanism. Technical mechanism used to ensure that the security policy is not violated by an adversary within the threat model, **we can only prepare for threats we're aware of**

(*Policy. ensure messages cannot be read by anyone but the sender and the receiver, Mechanism. encrypt the message before sending*)

Composition of Security Mechanisms

- **Defence in depth.** As long as one remains unbroken the Security Policy isn't broken (*two-factor auth*)

- **Weakest Link.** if anyone fails the Security Policy, it is broken. (*security questions for a lost password → just need to know the answer...*)

Humans can be vulnerabilities - phishing attacks, bad use of passwords...

To show a system is secure. (under a specific threat model)

Attacker - Just one way to violate **one** security property is enough.

Defender - No adversary strategy can violate the security policy.

Security Argument. Rigorous argument that security mechanisms in place are effective in maintaining security policy subject to assumptions on Threat Model.

5. **Separation of privilege.** No single accident, deception, or breach of trust is sufficient to compromise the protected information

- **Privilege.** A privilege allows a user to perform an action on a computer system that may have security consequences. (*create a file in a directory...*)

6. **Least Privilege.** Every program and every user of the system should operate using the least set of privileges necessary to complete the job. *Rights are added on need, discarded after use. Users should only know about things if they have to.*

7. **Least Common Mechanism** Minimize the amount of mechanism common to more than one user and depended on by all users. Every shared mechanism represents a potential information path between users.

8. **Psychological acceptability.** It is essential that the human interface be designed for ease of use, so that users routinely and automatically apply the protection mechanisms correctly. (*hide complexity, cultural acceptability...*)

9. **Work Factor**

Compare the cost of breaking the mechanism with the resources of a potential attacker. (*cost of compromising insiders, cost of finding a bug, monetization...*)

10. **Compromise recording** Detect and record security breaches with tamper-evident logs, ensuring traceability, integrity, confidentiality, and availability of recorded data.

User & Group Identities.

Most modern systems rely on DAC.

UIDs / GIDs. numerical identifiers for users and groups.

/etc/passwd: username : password :UID :GID :info :shell

/etc/group: defines secondary groups.

Each user has a home directory and belongs to one or more groups.

UNIX Model. Everything is a file.

directories	owner	group	others	owner	group	links	size	last modified	filename
drwxrwxr-x	1	catronco	catronco	4096	Sep 16 14:23	exampledir			
-rwxrwxrwx	1	catronco	catronco	8600	Sep 15 15:20	hello			
-rw-rw-rw-	1	catronco	catronco	150	Sep 15 15:14	hello.c			
-rw-rw----	1	catronco	catronco	45	Sep 15 15:07	test1.txt			

Directories. Read → list files; Write → add/remove files; Exec → traverse (cd).

Permission check order: 1. If process UID = file owner → check owner bits.

2. Else if GID matches → check group bits.

3. Else → check "other" bits.

Root (UID 0). bypasses all checks.

Changing Permissions.

chmod. modify permission bits. chmod +r file, chmod 666 file.

chown. change file owner/group(opt.) chown root:staff /srv/config

chgrp. change file group. chgrp www-data /var/www

Special Rights.

suid (set user ID). run program with owner's privileges. puts s in owner's x field chmod u+s filename

allows normal users to change passwords without full root access.

sgid (set group ID). run program with group's privileges.

on a directory - creating a file inside gets folder's group

d rwx r-s r-x 2 root staff /srv/shared

sticky bit. on directories, prevents deleting/renaming files you don't own.

d rwx rwx rwt 10 root tmp /tmp

Special Users.

root: UID 0, full privileges, bypasses checks, in TCB.

sudo run as root su [username] switch to account(default is root)

nobody: UID -2, owns no files, minimal privileges, safer for untrusted code.(least privilege)

⊕ flexible, simple model, widely adopted.

⊖ relies on **ambient authority**, prone to Confused Deputy attacks.

Windows: DACL.

Controls access to objects via list of ACEs.

ACE (Access Control Entry). <Type, Principal, Permissions, Flags>

- **Types.** Allow / Deny (negative / positive). Deny takes precedence and ordered with Denied permissions first.
- **Principal.** User or group (SID).
- **Permissions.** Fine-grained rights(Read, Write, Execute, Delete...)
- **Flags.** Inheritance, propagation, object-specific.

Access Tokens. Thread/process carries user + group SIDs checked against DACL.

Least Privilege. Users run limited by default; elevation via "Run as admin."

Capabilities.

associate permissions to subjects

Alice: {file1, read/write} Bob: {file2, read/write}, (file3, read/write)}

⊕ easy to determine all permissions of a user and to delegate rights by subject.

⊖ hard to determine who can access a resource or to revoke rights by resource.

Security Model. a design pattern for a set of properties. (\triangle not covered by model - who are the subjects? what are the objects? what mechanisms to use to implement it?)

Bell-La Padula (BLP).

security model where Subjects S and objects O are associated to a level of **Confidentiality**.

Access rights. Execute, Read, Append, Write.

- Objects are associated to a Security Level = (Classification, set of categories).

{(Unclassified < Confidential < Secret < Top Secret), {NATO, Crypto, Nuclear}}

Dominance Relationship.

Transitive. There always is a Top and Bottom, Only partial ordering (some pairs of elements can't be compared).

A security level (l_1, c_1) dominates (l_2, c_2) if and only if $l_2 \leq l_1$ and $c_2 \subseteq c_1$.

eg. $(\text{Secret}, \{\text{Nuclear, Army}\}) \geq (\text{Confidential}, \{\text{Army}\})$

Clearance. max security level a subject has been assigned. clearance-level(S_i).

Current Security level. subjects can operate at lower security levels. current-level(S_i).

Declassification. Removing classification labels from information so it can be shared with lower security levels. soldier (S) with Secret clearance shares sanitized info like "We've seen tanks moving with a lower-level user (U).

Covert Channels. Any channel that allows information flows contrary to the security policy, Storage channels(shared counters, ...), Timing channels(queueing time...).

\triangle Least Common Mechanism. mitigated by adding noise or isolation(no high \leftrightarrow low level communication).

BIBA (Integrity).

Security model where subjects S and objects O are 1. associated to a level of **Integrity** (trustworthiness).

Access rights. Read (Observe), Write (Modify)

BIBA to create Integrity Policies.

Simple Integrity Property.

If (s, o, r) is a current access, then level(s) dominates level(o).

(SUBJECTS CAN'T READ DOWN) \Rightarrow prevents contamination from low to high.

*-Integrity Property.

If (s, o, w) is a current access, then level(o) does not dominate level(s).

(SUBJECTS CAN'T WRITE UP) \Rightarrow prevents low

contamination from corrupting high state.

Discretionary Property (DAC). Subject must still have explicit permission for the access (r, w, x) per the access control matrix.

Basic Integrity Theorem (induction). If the initial state is integrity-secure and all state transitions satisfy BIBA properties, then all reachable states preserve integrity.

Sanitization (lifting low \rightarrow high).

Fail-safe default: deny by default; elevate only after checks pass.

- Whitelist over blacklist: validate that *all* required properties of "good" hold.

- Context-aware: encoding, schema, range, semantics, and provenance checks.

Note: Sanitization bugs commonly break integrity guarantees.

Principles for Integrity.

- Separation of Duties. Require multiple principals to perform an operation

- Rotation of Duties. Allow a principal only a limited time on any particular role and limit other actions while in this role

- Secure Logging. Tamper evident log to recover from integrity failures. Consistency of log across multiple entities is key.

Chinese Wall Model.

All objects are associated with a label denoting their origin.

(Pepsi, Coca-Cola, Microsoft Audit, Microsoft Investments)

Define conflict sets of labels. (Pepsi, Coca-Cola, {Microsoft Audit, Microsoft Investments}).

Subjects are associated with a history of their accesses to objects and their labels.

Access Rules. A subject can access an object (read or write) **only if** the access does not allow information flow between items with labels in the same conflict set.

Example (Direct Flow).

Example (Indirect Flow).

1. Access to Pepsi (OK)

2. Access to Microsoft Invest (OK)

3. Access to Coca-Cola (Denied)

Sanitization.

Allows more flexibility by "un-labeling" some items—controlled sharing/reuse of data.

All together (Asymmetric Cryptography) - Bob sends a message to Alice

1. Get public keys (via PKI).

Bob retrieves Alice's encryption public key PK_{Alice} and verification public key PK_{Alice} .

2. Prepare the message.

Bob has message M and computes its hash $h(M)$. Allows: the signature to cover the exact content of M while keeping computation efficient.

3. Sign the message hash.

Bob uses his secret signing key SK_{Bob} to generate a digital signature. $\text{Sign}_{SK_{Bob}}(h(M))$.

4. Encrypt the message.

Bob encrypts the message using Alice's encryption public key PK_{Alice} : $E_{PK_{Alice}}(M)$.

Allows: only Alice, who holds the matching secret key SK_{Alice} , to read the message.

Authenticity. only sender can generate a valid signature, allowing the server to verify the sender's identity and preventing impersonation. **Integrity.** Any message change invalidates signature, ensuring tampering is detectable. **Non-repudiation.** only sender can produce the valid signature \rightarrow they can't deny later having sent the message.

Block Cipher - Electronic Code Block (ECB). encrypt each message block independently using the same key.

1. The message M is divided into blocks.

$M = M_1 M_2 M_3 M_4$.

2. Each block is encrypted separately $C_i = E_K(C_i)$.

3. Decryption. $M_i = D_K(C_i)$.

\triangle If $(M_i = M_j) \Rightarrow (C_i = C_j) \Rightarrow$ (freq. analysis)

Block Cipher - Cipher Block Chaining (CBC)

add randomness/inter-block dependence by chaining blocks.

$C_0 = IV$ **Encryption.** $C_i = E_K(M_i \oplus C_{i-1})$

Decryption. $M_i = D_K(C_i) \oplus C_{i-1}$

IV must be random / isn't a ciphertext block but needed for decryption.

- If IV incorrect, only first plaintext block will decrypt incorrectly, later blocks remain unaffected.

Cryptography.

Data in transit. Securing communications. **Data at rest.** Securing stored informations Let C the Ciphertext, K the Key, M the Message/Plaintext. $C = E_K(M)$ and $M = D_K(C)$

Keyspace. Number of possible keys that can be used in an encryption algorithm.

Invertibility Requirement. $\forall K, M|D_K(E_K(M)) = M$ (otherwise can't recover message)

Security Requirement. Functions should be hard to invert without knowing K.

Ideal Case - adversary must try every possible combination of keys (brute-force).

Caesar's Cipher

- Encryption. Shift each letter by a fixed number (K)

- Decryption. Shift each letter by a fixed number (-K)

Keypspace. Only 25 possible keys.

$\log_2(25) = 4.6$ bits of security (too small).

Both can be broken using Frequency Analysis Attack.

Frequency Analysis. Use statistical properties of the language.

1. Most frequent letter in English is 'e'. 2. Identify most frequent letter in Ciphertext. 3. Map it to 'e'.

Ideally. An N-bit key should offer security as close to N bits as possible (require 2^n attempts), if not the algorithm is considered broken.

Types of Adversaries. security models

- Passive Eavesdropper - adversary can only read the ciphertext

- Active Attacker - adversary can influence the system (corruption of one of parties, ...)

Known Plaintext Attack (KPA). **Chosen Plaintext Attack (CPA).**

Active security model. Attacker gets - Suppose Eve convinces Alice to encrypt chosen messages

access to some pairs, (message, cipher) with the secret key text, all encrypted with secret K. - For all m , Eve gets access to $E_K(m)$. Eve has access to an En-

From these pairs, she tries to guess key crypton Oracle.

K to decrypt other messages

(realistic, she could get access to an encrypted messaging app realistic because of message headers, for limited time, or to an encryption api).

message headers ("From:...", times- \triangle if Eve encrypts entire alphahabet she can reveal the key instantly.

Side-Channels Attacks

- Timing attack - Eve measure how long it takes for a given message to get encrypted or a ciphertext to be decrypted

- Power Analysis - Eve observes the energy consumed by the device doing the crypto.

One-Time-Pad (OTP).

Goal - Remove frequency analysis

- Use a key of random bits as along as the message. - Equally likely

- Key K is uniformly random, so C gives the adversary zero new information about M.

Perfect Secrecy.

OTP is Perfectly Secure. **OTP is Perfectly Secure.**

- For any ciphertext C, every possible plaintext M is equally likely

- Eve should be random for every sent message. $\forall m, c, P(M = m|E_K(m) = c) = P(M = m)$.

Key should be random for every sent message. $\forall m, c, P(M = m|E_K(m) = c) = P(M = m)$.

Otherwise attacker can collect information about the message.

Integrity Attack Example (Eve flips the first bit of M): Reuse of key Attack. Send two messages of length 5 with key reuse.

1. Eve flips the first of C to get C'

2. Bob decryts C': $M' = C' \oplus K = (M \oplus K \oplus \Delta) \oplus K = M \oplus \Delta$

Symmetric Cryptography Schemes **Symmetric Cryptographic key**

Encryption/decryption done with the same key

Known to both parties. Partners must agree on the key before starting using the primitive

Block Ciphers. Operate on fixed-size blocks.

Stream Ciphers. Operate on bit/byte at a time, like pseudo-OTP.

Must be secret. Revealing the key eliminates any protection provided by the primitive

Stream Ciphers - The pseudo-OTP. emulate OTP while solving key-length problem.

1. Secret short Key (K) shared between Alice and Bob 2. Key Stream Generator - Uses K and Initialization Vector (IV) \rightarrow an arbitrary long, pseudo-random bit stream (S). 3. Encryption. $C = M \oplus S$ (generator needs a key as main seed to generate a predictable random sequence, an iv for the sequence to start differently on each run)

\triangle Speed, Low Error propagation(errors in one bit do not affect subsequent symbols)

\triangle Low Diffusion (a change in a bit only affects one bit), Susceptibility to Modification (low diffusion makes it easier to tamper)

\triangle key stream generators are periodic because seed is finite. Thus, we need period long enough to not be an issue (avoiding frequency analysis)

Linear Feedback Shift Register for Key Stream Generators. build a ""random"" sequence of bits using a linear recurrence relation on a sequence of bit.

Example: Starting with state $a_0, a_1, a_2, a_3 | a_n = a_{n-3} \oplus a_{n-4}$ (\oplus Easy to build/analyze.)

Distribution property.

if characteristic polynomial of the recurrence relation of the LFSR is primitive. The maximum possible number of states before zero state appears exactly once in a cycle.

repeating. For an L-bit register, the maximum period is $2^L - 1$ states

Symmetric ciphers are fast, but require a pre-shared secret key

Diffe-Hellman Setup. Alice, Bob (and Eve) know large public parameters: prime p/generator g

Alice's Actions

1. Chooses private secret a. 2. Computes Public Value $A = g^a \pmod p$ 3. Sends A to Bob.

Eve sees A, B, g, and p, but can't find a/b due to the Distcrete Log Problem.

Both Alice and Bob can compute a shared secret K: $K = B^a \pmod p = (g^b)^a \pmod p = A^b \pmod p$

this is uses Trapdoor functions, easy to compute but hard to revert.

Man-In-The-Middle (MITM). 1. Eve sends E_B to Alice. Alice computes K_{AE}

Vanilla DH is Vulnerable to Active Attack, 2. Eve sends E_A . Bob computes K_{BE}

DH guarantees key agreement, not authenticity. 3. Creates an Alice-Eve channel and Eve-Bob channel

Attacker could intercept A from Alice and B she can read, modify, and relay all communication from Bob, and then impersonate.

Asymmetric Cryptography. Encryption using a public-private key pair. Public key can be stored on a public server, the **Public Key Infrastructure (PKI)**.

Hash Functions. maps any-length message to a fixed-size output using a hash function. $h(M) = H$

Security properties. 1. Pre-image resistance. Given $H = h(M)$, it is hard to find M .

2. Second pre-image resistance. Given M , it is hard to find another $M' \neq M$ such that $h(M') = h(M)$.

3. Collision resistance. It is hard to find any two values M and M' such that $h(M) = h(M')$.

Block Cipher - Counter Mode **Block Cipher - CBC-MAC.** turns block cipher into MAC by chain-

add randomness per-block without ing message blocks with encryption.

inter-block dependencies. $C_0 = IV$ **Encryption.** $C_i = E_K(M_i \oplus C_{i-1})$ **MAC** $_K(M_1, \dots, M_n) = C_n$ (parallel)

Decrypt. $C_i = M_i \oplus E_K(\text{Nonce} + i)$. \triangle message length must be known and fixed; otherwise, extension

Decrypt. $M_i = C_i \oplus E_K(\text{Nonce} + i)$. attacks are possible (e.g., using $M(T \oplus M')$).

Nonce must be unique (never with the same key) reuse leaks keystream.

1. Encrypt-and-MAC. send $E_K(P_1 P_2)$ and $MAC_K(P_1 P_2)$. protects

plaintext only, not ciphertext; inefficient and leaks repetition.

2. MAC-then-Encrypt send $E_K(P_1 P_2 || MAC_K(P_1 P_2))$. hides MAC (MAC).

short tag computed from and provides confidentiality but ciphertext can be altered under a message and secret key(pre-tected).

3. Encrypt-then-MAC send $E_K(P_1 P_2 || MAC_K(E_K(P_1 P_2)))$. au-

thenticates ciphertext, ensures confidentiality/integrity ✓

shared) to verify message integrity

3. Encrypt-then-MAC send $E_K(P_1 P_2 || MAC_K(E_K(P_1 P_2)))$. au-

thenticates ciphertext, ensures confidentiality/integrity ✓