

Data Intensive Systems

IN BA6

Notes by Ali EL AZDI

February 16, 2026

Introduction

This document is designed to offer a concise, LaTeX-styled overview of the Data Intensive Systems course, with a focus on essential content and clarity. The **ConceptName.** notation is usually used for definitions in that same spirit. The latest version I am currently working on (not necessarily stable) is available at <https://github.com/elazdi-al/notes>. I do make mistakes often, and sharing this document is also a way to have my notes proofread. If you notice anything weird, incorrect, or missing, please do not hesitate to contact me at ali.elazdi@epfl.ch. I would be happy to answer and explain why I made certain choices.

with this being said, welcome back everyone and good luck!

Contents

Contents	3
1 L1 - Introduction	6
1.1 Why data management?	6
1.1.1 What is data?	6
1.1.2 What is a database?	7
1.1.3 What is a Database Management System (DBMS)?	7
1.1.4 What does a DBMS do?	7
1.1.5 Is your file system a DBMS?	7
1.1.6 Is the web a DBMS?	8
1.1.7 The scope of DBMS	8
1.1.8 What is the intellectual contribution?	9
1.2 DBMS Architecture	9
1.2.1 Describing data	9
1.2.2 Example: Schema of a University database	9
1.2.3 Describing data: Levels of abstraction	9
1.2.4 DBMS cares about data independence types	10
1.2.5 Simplified DBMS architecture	10
1.2.6 Relational data model	10
1.2.7 Database design requirements	11
1.2.8 What is a conceptual design	11
1.2.9 Conceptual Design: Entity-Relationship Model	11
1.2.10 Key constraints	13
1.2.11 Participation constraints	13
1.2.12 Weak entities	14
1.2.13 ISA (“is a”) hierarchies	14
1.2.14 Aggregation	15
1.2.15 Conceptual design using the ER model	16
2 L2 - Relational Model & SQL Basics	17
2.1 Relational model	17
2.1.1 Core concepts	17
2.1.2 Schema vs. instance	17
2.1.3 Null values	18

2.2	Keys and integrity constraints	18
2.2.1	Key hierarchy	18
2.2.2	Integrity constraints	18
2.3	SQL: Structured Query Language	19
2.3.1	DDL and DML	19
2.3.2	Most common SQL commands	19
2.3.3	Creating relations	19
2.3.4	Adding and deleting tuples	19
2.3.5	Keys in SQL	19
2.3.6	Foreign keys and referential integrity	20
2.3.7	Integrity constraints	20
2.4	Relational query languages	21
2.4.1	Formal relational query languages	21
2.4.2	Preliminaries	21
2.4.3	Example schema and instances	21
2.4.4	Five basic operations	22
2.4.5	Selection (σ)	22
2.4.6	Projection (π)	23
2.4.7	Composing operators	23
2.4.8	Rename operator (ρ)	24
2.4.9	Union operator ($\cup$)	24
2.4.10	Set-difference operator ($-$)	24
2.4.11	Compound operators	25
2.4.12	Cross-product operator ($\times$)	25
2.4.13	Join operator ($\bowtie$)	25
2.4.14	Complex queries	26
2.4.15	Division operator ($/$)	27
3	SQL-1 - SQL Introduction	28
3.1	Relational terminology	28
3.2	Basic SQL queries	28
3.2.1	The simplest query	28
3.2.2	Showing specific columns (Projection)	29
3.2.3	Showing specific rows (Selection)	29
3.2.4	SQL vocabulary	29
3.2.5	Evaluation order	29
3.3	Querying multiple relations	29
3.3.1	Naïve evaluation in steps	30
3.4	Aggregate operators	30
3.4.1	Examples	31
3.4.2	Name and age of the oldest sailor(s)	31
3.5	GROUP BY	31
3.5.1	Using GROUP BY	32

3.5.2	HAVING — filtering groups	32
-------	-------------------------------------	----

Chapter 1

L1 - Introduction

1.1 Why data management?

Modern software is data-driven: systems are useful only when data can be stored, shared, and queried correctly over time. This chapter builds that idea step by step, from raw data to DBMS architecture.

1.1.1 What is data?

We first separate *raw facts* from *usable meaning*.

Data. Raw facts used as input for reasoning, discussion, and computation. Data can be useful, irrelevant, or redundant. It must be processed to become meaningful.

Information. Processed data with meaning in context. Information is relevant to a specific problem and actionable for decision-making.

Practical flow:

$$\text{Data} \xrightarrow{\text{organize} + \text{process}} \text{Information}$$

Information then supports higher-level outcomes such as **knowledge** and **wisdom**.

1.1.2 What is a database?

Once data has meaning, we need structure.

Database. A large, integrated, structured collection of data, usually modeling a real-world enterprise.

Entity. An object type we store data about (for example, a student or a course).

Relationship. A meaningful association between entities (for example, a student enrolls in a course).

Example: University. Entities are courses, students, and professors; relationships include enrollment and teaching.

1.1.3 What is a Database Management System (DBMS)?

A database alone is passive. We need software that manages access and correctness.

DBMS. A software system designed to store, manage, and facilitate access to databases.

A useful decomposition is:

$$\text{DBMS} = \text{interrelated data (database)} + \text{programs to access/manage it}$$

1.1.4 What does a DBMS do?

Core role of a DBMS. Provide efficient, reliable, convenient, and safe multi-user storage of and access to massive persistent data.

In practice, this means:

- failure protection against hardware/software faults, power loss, and malicious behavior,
- high throughput (thousands of queries/updates per second),
- high availability (24/7 operation),
- concurrency control for correct multi-user access,
- support for very large-scale data,
- persistence (data outlives programs),
- high-level declarative access with physical data independence.

Persistent data. Data that remains stored beyond the lifetime of a process or program execution.

Concurrency control. Mechanisms that coordinate simultaneous operations so outcomes remain correct.

1.1.5 Is your file system a DBMS?

This thought experiment shows why plain files are often insufficient for data-intensive multi-user applications.

Thought experiment 1: concurrent edits

Two collaborators edit and save the same file at nearly the same time. The natural question is: *whose changes survive?* Outcomes can be yours, your partner's, both, neither, or simply **unclear**.

Thought experiment 2: update during power loss

A file update is in progress and power fails. Which changes survive: all, none, only changes since last save, or again **unclear**?

The key issue is the missing contract: guarantees are unclear, which makes application logic fragile and hard to reason about.

How do you write programs over a platform that promises “???”?

If correctness, durability, isolation, and recovery guarantees are not explicit, every application must re-implement them. **A DBMS exists to provide these guarantees systematically.**

1.1.6 Is the web a DBMS?

The web offers strong information retrieval, but retrieval is not equivalent to database management.

The web is not a DBMS

- data is mostly unstructured and untyped,
- the “correct” answer for many queries is not well-defined,
- completeness is not guaranteed,
- data manipulation guarantees are limited,
- freshness, cross-item consistency, and fault-tolerance guarantees are weaker or undefined.

Practical architecture

Most serious websites place a DBMS behind the web layer (for example, `facebook.com` with MySQL).

Is ChatGPT a DBMS???

Short answer. No. ChatGPT is a language model interface, not a DBMS.

It can connect to database-backed systems, but by itself it does not provide the full DBMS contract (schema-centered storage, transaction guarantees, and complete declarative data management).

1.1.7 The scope of DBMS

The question is no longer only “*where to store bytes*”; it is how to query, update, and trust data under real workload conditions.

What more could we want than a file system?

- simple and efficient *ad-hoc* querying,
- robust concurrency control,
- recovery mechanisms,
- the benefits of well-designed data models.

Ad-hoc query. A query formed for a specific immediate need, not predefined as part of a fixed workflow. *With many queries, users, transactions, and resource constraints, complexity increases quickly. Managing that complexity is a core DBMS concern.*

Can the OS offer services like memory management etc? Again, not really

The operating system is essential, but it does not provide the complete data-management contract that applications need.

1.1.8 What is the intellectual contribution?

DBMS is not only engineering practice; it also contributes key ideas about representation, semantics, and correctness.

Core contributions of DBMS thinking

- **Representation of information:** data modeling,
- **Languages and systems for querying data:** complex queries with real semantics over massive data,
- **Concurrency control for data manipulation:** controlled concurrent access with transactional meaning,
- **Reliable data storage:** preserving semantics even after failures (for example, power loss).

Transactional semantics. The guarantee that concurrent and failing executions still behave like correct transactions under clear correctness rules.

1.2 DBMS Architecture

The main architectural question is simple: how do we keep one coherent meaning of data while serving different users and storing bytes efficiently on disk?

1.2.1 Describing data

Data model. A collection of concepts for describing data.

Data models are high-level abstractions: they hide low-level storage details and let us reason about structure and meaning first. Typical model families include *relational*, *hierarchical*, *graph*, and *object-oriented*.

Relation. A named table with rows (tuples) and columns (attributes).

Schema. The structural definition of a relation: attribute names, types, and organization.

In the **relational model**, data is represented as a set of records organized into relations. In **nested models**, not all data is forced into flat tables; hierarchies and arrays are natural first-class structures.

Schema vs. data. Schema is the type (for example, *string*); data is a value instance (for example, "hello") that follows that type.

1.2.2 Example: Schema of a University database

We use one running example throughout the rest of the section.

Students `sid:string, name:string, login:string, age:integer, gpa:real`

Enrolled `sid:string, cid:string, grade:string`

Courses `cid:string, cname:string, credits:string`

Some later examples use `credits:integer`. The important point is explicit typing in the schema.

1.2.3 Describing data: Levels of abstraction

External schema. User- or application-specific views, often used for access control.

Logical schema (conceptual level). The main global structure seen by users and application programs.

Physical schema (internal level). How data is stored on disk: files, indexes, and storage layout choices.

The mapping is:

External schemas $\rightarrow$ Logical schema $\rightarrow$ Physical schema

Many external views map to one logical schema that is then implemented by physical structures.

Running example across levels

External views (access control)

Students_Info `sid:string, name:string`

Course_Enrollment `cid:string, #enrolled:integer`

Logical schema (what applications see)

The logical level uses the university relations above: **Students**, **Courses**, and **Enrolled** (with `credits` often shown as `integer` in this level).

Physical schema (how storage is organized)

Relations can be stored as unordered files, with indexes such as `Students.sid` and `Courses.cid`.

1.2.4 DBMS cares about data independence types

Data independence. The ability to change the schema at one level of a database system without changing the schema at the next higher level.

Logical data independence. The ability to change the conceptual schema without changing user views.

Physical data independence. The ability to change the internal schema without changing the conceptual schema or user views.

This matters because data-intensive systems evolve constantly: applications should keep working while storage layouts, indexes, and conceptual structures improve over time.

1.2.5 Simplified DBMS architecture

The architecture has two linked flows: a *design flow* to define data and a *query flow* to access data efficiently.

Design flow (from requirements to storage)

- **Conceptual design:** capture the enterprise at a high level, typically with ER models.
- **Logical design:** translate ER structures into a relational model.
- **Physical design:** choose indexing and storage structures.
- **Database storage:** persist data in disk files.

Query flow (from SQL to results)

Applications issue SQL (rooted in relational algebra). Inside the DBMS, processing is layered:

- query optimization and execution,
- relational operators,
- access methods,
- buffer management,
- disk space management.

Cross-cutting services enforce correctness and performance across these layers: query processing, transaction management, and concurrency control.

1.2.6 Relational data model

In practice, relational design relies on rows/columns and explicit links between relations through keys.

Key. An attribute (or set of attributes) whose value identifies a tuple uniquely inside a relation.

Foreign key. An attribute (or set) in one relation whose values reference keys in another relation, linking the two relations.

Example relation instance

	SID	CID	Grade		SID	Name	Login	Age	GPA
Enrolled	53666	Carnatic101	5	Students	53666	Jones	jones@cs	18	5.4
	53666	Raggae203	5.5		53688	Smith	smith@eecs	18	4.2
	53650	Topology112	6		53650	Smith	smith@math	19	4.8
	53666	History105	5						

Here, `Enrolled.SID` acts as a foreign key referencing `Students.SID`.

This raises the next design question: *how do we design such schemas systematically?*

1.2.7 Database design requirements

Database design proceeds in stages, from problem understanding to implementation choices.

- **Requirements analysis:** identify user needs and what the database must support.
- **Conceptual design:** produce a high-level enterprise description (often with ER).
- **Logical design:** translate ER into the DBMS data model.
- **Schema refinement:** improve consistency and normalization.
- **Physical design:** choose indexes and disk layout.
- **Security design:** define who can access what.

1.2.8 What is a conceptual design

Conceptual design. A high-level description of the enterprise data, independent of storage and implementation details.

It answers three foundational questions:

- what are the *entities* and *relationships* in the enterprise?
- what information about them should the database store?
- which *integrity constraints* must always hold?

ER diagram. A pictorial representation of the conceptual schema in the ER model.

An ER diagram can be mapped into a relational schema; this bridge connects conceptual understanding to implementable database design.

1.2.9 Conceptual Design: Entity-Relationship Model

The ER model is the standard language for conceptual design: it captures entities, attributes, relationships, and integrity constraints before logical and physical design decisions.

Basics of ER model

Entity. A real-world object, distinguishable from other objects.

Attribute. A property used to describe an entity in the database.


Entity set. A collection of similar entities that share the same attributes.

Domain. The set of allowed values for an attribute.

Key (in ER). An attribute (or minimal attribute set) that uniquely identifies entities in an entity set.

Entity-set schema example

For the entity set **Employees**, we use attributes **ssn**, **name**, and **lot**, with **ssn** as key.



Instance of an entity set. The currently stored collection of entities in that set.

Example instances for **Employees**:


- (123-22-3666, John, P1)
- (231-31-5368, Tom, P2)

Relationships in ER model

Relationship. An association among two or more entities.


Relationship set. A collection of similar relationships.

Relationships may have their own attributes. Example: **Works_In** links **Employees** and **Department**, with relationship attribute **since**.



Instance of a relationship set. The currently stored set of relationship tuples between participating entity sets.

Example instance of **Works_In(ssn, did, since)** with links to participating entity instances:




Role names and repeated participation

The same entity set can participate in different relationship sets, and can also appear multiple times in one relationship set under different roles.

Role name. A label identifying the function of an entity set when it appears more than once in the same relationship.


Example: in **Reports_To**, **Employees** appears twice with roles *supervisor* and *subordinate*; the same entity set also participates in other relationships such as **Works_In**.



1.2.10 Key constraints

Key constraints restrict how entities participate in relationships. An arrow from an entity set to a relationship set means each entity participates in *at most one* relationship instance.

Key constraint. A restriction stating that each entity in a set can participate in at most one relationship instance of a given relationship set.




- **Many-to-Many:** no arrows; an employee can work in many departments, and a department can have many employees.
- **One-to-Many:** arrow on one side; each department has at most one manager (key constraint on Manages).
- **One-to-One:** arrows on both sides; each driver drives at most one vehicle and each vehicle has at most one driver.

1.2.11 Participation constraints

Participation constraints specify whether *every* entity in a set must participate in a relationship set.

Total participation. Every entity in the set must participate in at least one relationship instance. Drawn with a **thick line**.

Partial participation. Some entities may not participate in any relationship instance. Drawn with a thin line (default).



Combining key and participation constraints gives precise cardinality semantics:

- thick line + no arrow = total, many: every entity participates, possibly in multiple instances.
- thick line + arrow = total, at most one: every entity participates in exactly one instance.
- thin line + arrow = partial, at most one: an entity participates in zero or one instance.
- thin line + no arrow = partial, many: an entity may or may not participate, with no upper bound.

1.2.12 Weak entities


Weak entity. An entity that can be identified uniquely only by considering the primary key of another (*owner*) entity together with its own partial key.

Identifying relationship. The relationship set that links a weak entity set to its owner entity set. The weak entity must have total participation in this relationship.

Rules for weak entities:

- Owner and weak entity set participate in a one-to-many relationship (one owner, many weak entities).
- The weak entity set must have **total participation** in the identifying relationship.
- Weak entities have only a **partial key** (drawn with a dashed underline): it uniquely identifies weak entities only within the scope of a given owner.

Notation: the weak entity and its identifying relationship are drawn with **thick borders**; the arrow from the weak entity to the identifying relationship indicates the key constraint.




Here, **Dependents** is the weak entity owned by **Employees** through the identifying relationship **Policy**. The attribute **pname** is a partial key.

1.2.13 ISA (“is a”) hierarchies

ISA hierarchy. A relationship between a superclass entity set and one or more subclass entity sets, where each subclass entity “is a” member of the superclass. Attributes are inherited from superclass to subclass.

If we declare A **ISA** B, then every A entity is also considered a B entity and inherits all attributes of B.



Hourly_Emps and **Contract_Emps** are subclasses of **Employees**: they inherit **ssn**, **name**, and **lot**, and add their own specialized attributes.

Two orthogonal design constraints govern ISA hierarchies:


- **Overlap constraint:** can an entity belong to more than one subclass? Default: *no* (disjoint).
- **Covering constraint:** must every superclass entity belong to at least one subclass? Default: *no* (partial coverage).

Reasons to use ISA:

- add descriptive attributes specific to a subclass,
- identify entity sets that participate in specific relationship sets.

ISA example: Student hierarchy

A student can be classified by degree type (Bachelor, Post-Graduate) *and* by enrollment status (Full-time, Part-time) simultaneously. These are two **overlapping** ISA branches from the same superclass. Within each branch, the subclasses are **non-overlapping** (a student is either a Bachelor or a Post-Graduate, not both).




Post-Graduate further specializes into Master and Doctorate (non-overlapping). Each subclass can participate in its own relationship sets: Bachelor students register for Projects; Post-Graduate students register for a Thesis.

1.2.14 Aggregation

Aggregation. A mechanism that treats a relationship set and its participating entity sets as a single higher-level entity set, so that it can participate in other relationships.

Aggregation is needed when we want to model a relationship *involving a relationship set*. For example: departments sponsor projects (the relationship **Sponsors**), and an employee needs to monitor each sponsorship.

Without aggregation we would need a ternary relationship among Projects, Departments, and Employees, which does not capture that monitoring applies to the *sponsorship* itself. With aggregation, the Sponsors relationship is enclosed in a dashed box and treated as an entity set that participates in **Monitors**.



1.2.15 Conceptual design using the ER model

Designing an ER schema involves several recurring choices. The right answer depends on the semantics of the data and the queries the application must support.

Design choices

- Should a concept be modeled as an *entity* or an *attribute*?
- Should a concept be modeled as an *entity* or a *relationship*?
- Should relationships be *binary* or *ternary*? Or use *aggregation*?

Entity vs. attribute

Should *address* be an attribute of Employees or an entity related to Employees?

It depends on how we use the data:

- If there are **several addresses per employee**, address must be an entity (attributes cannot be set-valued in the basic ER model).
- If the **structure** of an address matters (city, street, zip), address must be an entity (attribute values are atomic).

The same reasoning applies to relationship attributes. For example, if **Works_In** has attributes **from** and **to** to record time periods, an employee can work in a department for only one period. If we want to record *several* periods, we must promote the time information into a separate entity (**Duration**) linked by a new relationship.

Entity vs. relationship

Should a concept be a relationship or an entity with relationships to the things it connects? The guideline is: if the concept has **its own identity** or participates in **other relationships**, make it an entity.

Binary vs. ternary relationships

If a ternary relationship can be decomposed into binary relationships without loss of information, prefer binary. If not (some combinations of pairs do not imply the triple), a ternary relationship is needed.

Constraints in the ER model

- A lot of data semantics *can* (and should) be captured in ER diagrams: key constraints, participation constraints, weak entities, ISA hierarchies, and aggregation.
- However, some constraints *cannot* be expressed in ER diagrams (for example, “a manager must also be an employee of the managed department”). These must be enforced by the application or by database-level integrity constraints.

Chapter 2

L2 - Relational Model & SQL Basics

2.1 Relational model

The relational model is the foundation of modern database systems. It represents all data as tables and provides a clean, mathematically grounded framework for querying and constraining data.

2.1.1 Core concepts

A **database** is a set of named *relations* (*named tables with rows and columns*). Each attribute carries a **type** (domain): integer, real, string, file formats, enumerated, etc.

Tuple. A single row in a relation, containing one value for each attribute. All tuples in a relation are distinct and have no inherent order.

Running example:

Students

sid	name	login	age	gpa
50000	Dave	dave@cs	19	3.3
53666	Jones	jones@cs	18	3.4
53688	Smith	smit@ee	18	3.2
...	...	...	...	...

Colleges

name	location	strength
MIT	USA	10000
Oxford	UK	22000
EPFL	CH	9000
...	...	...

2.1.2 Schema vs. instance

A relation has two complementary aspects: its *schema* (*structural definition with attribute names and types*) and its content at a given moment.

Example schema: `Students(sid: string, name: string, login: string, age: integer, gpa: real)`.

Instance. The actual set of tuples stored in a relation at a given point in time.

Two measures describe an instance:

- **Cardinality:** the number of tuples (rows).
- **Arity** (or degree): the number of attributes (columns).

2.1.3 Null values

Not every attribute value is always known. A student who has not received any grades yet has no meaningful GPA.

Null value. A special marker representing an “unknown” or “undefined” attribute value. Nulls are not the same as zero or an empty string.

2.2 Keys and integrity constraints

Keys (attributes that uniquely identify tuples) enforce uniqueness and link relations together. Here we refine the hierarchy of key types and formalize integrity constraints.

2.2.1 Key hierarchy

Superkey. A set of attributes such that no two distinct tuples can share the same values in all those fields. Any superset of a key is also a superkey.

Candidate key. A *minimal* superkey: it is a superkey, and no proper subset of its fields is also a superkey.

Primary key. The candidate key chosen by the database administrator (DBA) as *the* key of the relation.

Example with **Students**(sid, name, login, age, gpa): both **sid** and **login** are candidate keys. The DBA picks **sid** as the primary key.

- Every superset of **sid** (e.g., {**sid**, **name**}, {**sid**, **login**, **gpa**}, ...) is a superkey but not a candidate key.
- **sid** alone and **login** alone are the only candidate keys.

2.2.2 Integrity constraints

A key is an **integrity constraint** (IC): a condition that must hold on every valid database instance.

Foreign keys (attributes referencing the primary key of another relation) establish **referential integrity**: every foreign-key value must match an existing key in the referenced relation.

2.3 SQL: Structured Query Language

SQL is the standard language for interacting with relational databases. It splits into two sub-languages.

2.3.1 DDL and DML

Data Definition Language (DDL). The subset of SQL used to create, modify, and delete relations, specify constraints, and administer users and security.

Data Manipulation Language (DML). The subset of SQL used to query, insert, update, and delete tuples.

2.3.2 Most common SQL commands

```
CREATE TABLE <name> (<field> <domain>, ...)
INSERT INTO <name> (<field names>) VALUES (<field values>)
DELETE FROM <name> WHERE <condition>
UPDATE <name> SET <field name> = <value> WHERE <condition>
SELECT <fields> FROM <name> WHERE <condition>
```

2.3.3 Creating relations

A **CREATE TABLE** statement defines a new relation by listing each attribute and its domain. The type of each field is enforced by the DBMS whenever tuples are added or modified.

Students – holds data about students:

```
CREATE TABLE Students
(sid CHAR(20),
 name CHAR(20),
 login CHAR(10),
 age INTEGER,
 gpa FLOAT)
```

Enrolled – holds course enrollments:

```
CREATE TABLE Enrolled
(sid CHAR(20),
 cid CHAR(20),
 grade CHAR(2))
```

2.3.4 Adding and deleting tuples

Insert a single tuple:

```
INSERT INTO Students (sid, name, login, age, gpa)
VALUES ('53688', 'Smith', 'smith@cs', 18, 3.2)
```

Delete all tuples satisfying a condition:

```
DELETE FROM Students S WHERE S.name = 'Smith'
```

2.3.5 Keys in SQL

The **PRIMARY KEY** clause declares the primary key; **UNIQUE** marks additional candidate keys.

```
CREATE TABLE Students
(sid CHAR(20),
 name CHAR(20),
 login CHAR(10),
 age INTEGER,
 gpa FLOAT,
 PRIMARY KEY(sid))
```

```
CREATE TABLE Person
(ssn CHAR(9),
 name CHAR(20),
 licence# CHAR(10),
 PRIMARY KEY(ssn),
 UNIQUE(licence#))
```

Choosing keys carefully

Keys must be used carefully: they encode real-world constraints. Consider **Enrolled**: “for a given student and course, there is a single grade.”

Correct – composite PK on (sid, cid):

```
CREATE TABLE Enrolled
(sid CHAR(20),
 cid CHAR(20),
 grade CHAR(2),
 PRIMARY KEY(sid, cid))
```

Wrong – PK on sid alone with UNIQUE(cid, grade) encodes a different constraint: “students can take only one course, and no two students in a course receive the same grade.”

2.3.6 Foreign keys and referential integrity

Foreign keys (fields referencing the primary key of another relation) act as logical pointers between relations. If all foreign-key constraints are enforced, the database achieves **referential integrity** (no dangling references).

Students

sid	name	login	age	gpa
50000	Dave	dave@cs	19	3.3
53666	Jones	jones@cs	18	3.4
53688	Smith	smit@ee	18	3.2
...	...	...	...	...

Enrolled

cid	sid	grade
Carnatic101	53666	C
Reggae203	50000	B
Topology112	53666	A
...	...	...

Here **Enrolled.sid** is a foreign key referencing **Students.sid**: every enrolled student must actually exist in **Students**.

Enforcing referential integrity

The DBMS must decide what to do when a foreign-key constraint is threatened:

- **Insert** an Enrolled tuple with a non-existent sid → reject the insert.
- **Delete** a Students tuple that is referenced by Enrolled → several policies:
 - cascade-delete all Enrolled tuples that refer to it,
 - disallow the deletion entirely,
 - set the referencing sid to a default value,
 - set it to NULL (denoting “unknown” or “inapplicable”).
- **Update** the primary key of a Students tuple → same policies apply.

2.3.7 Integrity constraints

Integrity constraint (IC). A condition that must be true for *any* instance of the database (e.g., domain constraints, key constraints, foreign-key constraints). ICs are specified when the schema is defined and checked whenever relations are modified.

Legal instance. An instance of a relation that satisfies all specified ICs. The DBMS should not allow illegal instances.

When the DBMS enforces ICs, stored data is more faithful to real-world meaning and data-entry errors are caught early.

2.4 Relational query languages

Query language (QL). A language that allows manipulation and retrieval of data from a database. Unlike programming languages, QLs are not Turing-complete: they are designed for easy, efficient access to large data sets, not general computation.

The relational model supports simple yet powerful QLs with a strong formal foundation based on logic, enabling significant optimization.

2.4.1 Formal relational query languages

Two mathematical query languages form the basis for real languages like SQL:

Relational Algebra. An *operational* language: it specifies *how* to compute a result step by step. Useful for representing execution plans inside database engines.

Relational Calculus. A *declarative* language: it specifies *what* the result should look like, not how to compute it. Non-procedural.

Codd's Theorem. Relational algebra and relational calculus are equivalent in expressive power: for every query in one, there is an equivalent query in the other.

Understanding both is key to understanding SQL and query processing.

2.4.2 Preliminaries

A query is applied to *relation instances*, and the result of a query is also a relation instance.

- The *schemas* of input relations are fixed (but the query runs over any legal instance).
- The schema of the *result* is also fixed, determined by the query language constructs.

Two notations are used:

- **Positional notation:** easier for formal definitions.
- **Named-field notation:** more readable; both are used in SQL.

2.4.3 Example schema and instances

We use a university database throughout the relational algebra examples:

instructor	<u>ID</u> , name, dept_name, salary
course	<u>course_id</u> , title, dept_name, credits
teaches	<u>ID</u> , <u>course_id</u> , <u>sec_id</u> , <u>semester</u> , <u>year</u>
section	<u>course_id</u> , <u>sec_id</u> , <u>semester</u> , <u>year</u> , building, room_number, time_slot_id
student	<u>ID</u> , name, dept_name, tot_cred
takes	<u>ID</u> , <u>course_id</u> , <u>sec_id</u> , semester, year, grade

Running example instance of `instructor`:

instructor

ID	name	dept_name	salary
22222	Einstein	Physics	95000
12121	Wu	Finance	90000
32343	El Said	History	60000
45565	Katz	Comp. Sci.	75000
98345	Kim	Elec. Eng.	80000
76766	Crick	Biology	72000
10101	Srinivasan	Comp. Sci.	65000
58583	Califieri	History	62000
83821	Brandt	Comp. Sci.	92000
15151	Mozart	Music	40000
33456	Gold	Physics	87000
76543	Singh	Finance	80000

2.4.4 Five basic operations

Relational algebra has five primitive operators, split into two groups:

Unary operators (one input relation):

- **Selection** (σ): output a subset of *rows* (horizontal filtering).
- **Projection** (π): output a subset of *columns* (vertical filtering).

Binary operators (two input relations):

- **Cross-product** ($\times$): concatenate every row of the first relation with every row of the second.
- **Set-difference** ($-$): output tuples in R that are not in S .
- **Union** ($\cup$): output tuples in R or in S (or both).

Since each operation returns a relation, **operations can be composed**.

2.4.5 Selection (σ)

Selection outputs only the rows that satisfy a given condition. The output schema is the same as the input.

$$\sigma_{\text{condition}}(\text{relation})$$

Without condition

The simplest expression $\sigma(\text{instructor})$ returns all rows unchanged (equivalent to `SELECT * FROM instructor`).

With a single condition

$$\sigma_{\text{dept_name}=\text{"Physics"}}(\text{instructor})$$

```
SELECT * FROM instructor WHERE dept_name = 'Physics'
```

ID	name	dept_name	salary
22222	Einstein	Physics	95000
33456	Gold	Physics	87000

With a compound condition

Conditions can be combined with $\wedge$ (and), $\vee$ (or), and $\neg$ (not):

$$\sigma_{dept_name = \text{"Physics"} \wedge salary > 90000}(instructor)$$

```
SELECT * FROM instructor WHERE dept_name = 'Physics' AND salary > 90000
```

ID	name	dept_name	salary
22222	Einstein	Physics	95000

2.4.6 Projection (π)

Projection retains only the attributes in the projection list. The output schema contains exactly those attributes.

$$\pi_{\text{attribute list}}(relation)$$

$$\pi_{ID, name, salary}(instructor)$$

corresponds to:

```
SELECT ID, name, salary FROM instructor
```

The result has three columns instead of four (`dept_name` is dropped), but all 12 rows are kept because each row is still distinct.

Duplicate elimination

In relational algebra, projection **eliminates duplicates** because the result must be a set of tuples.

$$\pi_{dept_name}(instructor)$$

reduces the 12 rows to 7 distinct department names. In SQL, duplicates are kept by default; to match the algebra, use `SELECT DISTINCT`:

```
SELECT DISTINCT dept_name FROM instructor
```

2.4.7 Composing operators

Since every operator returns a relation, the output of one can feed into another. Selection followed by projection gives specific columns from filtered rows:

$$\pi_{name}(\sigma_{dept_name = \text{"Physics"}}(instructor))$$

```
SELECT name FROM instructor WHERE dept_name = 'Physics'
```

name
Einstein
Gold

First σ filters to Physics instructors, then π keeps only the `name` column. The order matters: projecting first would discard the column needed for selection.

2.4.8 Rename operator (ρ)

The rename operator (ρ) renames a relation or its attributes. Since algebraic expressions can yield unnamed relations natively, renaming is essential to avoid naming conflicts or to self-reference.

- Renames the list of attributes in the form `oldname  $\rightarrow$  newname` or `position  $\rightarrow$  newname`.
- Can be used to rename the output relation itself entirely.
- **Output schema:** same as input except for the renamed attributes.
- **Output instances:** returns the exact same tuples as the input.

Renaming the relation itself to i :

$$\rho_i(\text{instructor})$$

Renaming the ID attribute to instructor.ID :

$$\rho_{ID \rightarrow \text{instructor.ID}}(\text{instructor})$$

In SQL, this is achieved using the `AS` keyword in the `SELECT` or `FROM` clauses:

```
SELECT i.name FROM instructor AS i WHERE i.ID = ...
```

2.4.9 Union operator ($\cup$)

The union operator takes two input relations, which **must** be *union-compatible*:

- They must have the same number of fields.
- The “corresponding” fields must have exactly the same type.

Like projection, union in strict relational algebra implies *duplicate elimination*.

$$c1 \cup c2$$

In SQL, `UNION` implicitly eliminates duplicates, while `UNION ALL` is faster but keeps them.

```
(SELECT * FROM c1)
UNION
(SELECT * FROM c2)
```

2.4.10 Set-difference operator ($-$)

The set-difference operator outputs the set of tuples in $c1$ which are not in $c2$.

- The two input relations must also be *union-compatible*.
- Set difference is **not** commutative ($c1 - c2 \neq c2 - c1$).

$$c1 - c2$$

In SQL, this corresponds to the `EXCEPT` clause:

```
(SELECT * FROM c1)
EXCEPT
(SELECT * FROM c2)
```

2.4.11 Compound operators

Alongside the five basic operators ($\sigma, \pi, \times, -, \cup$), there are several additional **compound operators**. These add **no computational power** to the language (they can all be expressed using the basic operators), but they serve as useful, compact shorthands.

Intersection ($\cap$)

Intersection outputs the overlapping tuples from two *union-compatible* relations.

$$R \cap S = R - (R - S)$$

```
(SELECT * FROM c1)
INTERSECT
(SELECT * FROM c2)
```

2.4.12 Cross-product operator ($\times$)

The Cartesian product $S \times R$ pairs every single row of S with every single row of R .

- **Result schema:** has one field per field of S and R , inheriting the field names linearly.
- **Naming conflicts:** Both S and R may have fields sharing the exact same name. In this case, we use the renaming operator (ρ) to disambiguate.

$$instructor \times teaches$$

2.4.13 Join operator ($\bowtie$)

A simple $instructor \times teaches$ associates *every* instructor tuple with *every* teaches tuple. Most of these rows hold information about an instructor paired with a course they never taught!

To get only tuples of instructors and the specific courses they actually taught, we apply a selection:

$$\sigma_{instructor.id=teaches.id}(instructor \times teaches)$$

Because this pattern (σ after $\times$) is overwhelmingly common in databases, it gets its own dedicated operator called a **Join**.

$$instructor \bowtie_{instructor.id=teaches.id} teaches$$

If the matched columns have the identical name in both tables, the condition can be completely omitted. This is called a **Natural Join**.

In SQL, a join is explicitly specified using **JOIN**:

```
SELECT * FROM instructor
JOIN teaches ON instructor.ID = teaches.ID
```

2.4.14 Complex queries

Operators can be chained to answer sophisticated questions. Consider: “Find the names of all instructors in the Music department together with the course title of all the courses that they teach.”

This query involves three relations (`instructor`, `teaches`, `course`) and requires joining them before projecting the desired columns.

Two equivalent formulations:

(a) Apply the selection *after* the full join:

$$\pi_{name, title}(\sigma_{dept_name=\text{“Music”}}(instructor \bowtie (teaches \bowtie \pi_{course_id, title}(course))))$$


(b) Push the selection *before* the join:

$$\pi_{name, title}((\sigma_{dept_name=\text{“Music”}}(instructor)) \bowtie (teaches \bowtie \pi_{course_id, title}(course)))$$


Both return the same result, but formulation (b) is significantly more efficient: the selection filters the `instructor` relation *before* the expensive join, drastically reducing the number of intermediate tuples.

Query optimization

This illustrates a key principle: **query optimization determines performance**. The DBMS query optimizer automatically rewrites expressions (e.g. pushing selections down) to find an efficient execution plan. The two expression trees below show the difference visually:



(a) Initial: σ applied *after* joining all three relations — processes every tuple.



(b) Optimized: σ pushed *down* to filter `instructor` before the join — far fewer tuples flow upward.

Both trees compute the same result, but (b) is far more efficient because the expensive join operates on a tiny filtered relation instead of the full table.

2.4.15 Division operator (/)

Division is a compound operator useful for expressing “for all” queries (e.g., “find all suppliers who supply *every* part”).

Division (A/B). Given relations $A(x, y)$ and $B(y)$, the division A/B returns all values of x such that for **every** tuple b in B , the tuple (x, b) exists in A .

Worked example

Consider relation $A(\text{sno}, \text{pno})$ and three divisors $B_1 \subseteq B_2 \subseteq B_3$:

$$B_1 = \{p2\} \quad B_2 = \{p2, p4\} \quad B_3 = \{p1, p2, p4\}$$

$$A/B_1 = \{s1, s2, s3, s4\} \quad A/B_2 = \{s1, s4\} \quad A/B_3 = \{s1\}$$

A

sno	pno
s1	p1
s1	p2
s1	p3
s1	p4
s2	p1
s2	p2
s3	p2
s4	p2
s4	p4

As B grows (more conditions to satisfy), the result of A/B shrinks: only suppliers matching *every* part in B survive.

Expressing division with basic operators

Division can be expressed using only π , $\times$, and $-$.

Idea: find all x values that are *not* “disqualified”. An x is disqualified if pairing it with some $y \in B$ yields a tuple *not* in A .

Disqualified x values:

$$\pi_x((\pi_x(A) \times B) - A)$$

Final result:

$$A/B = \pi_x(A) - \pi_x((\pi_x(A) \times B) - A)$$

Step by step with $B_3 = \{p1, p2, p4\}$

- $\pi_{sno}(A) = \{s1, s2, s3, s4\}$
- $\pi_{sno}(A) \times B_3$ produces all $4 \times 3 = 12$ possible (sno, pno) pairs.
- $(\pi_{sno}(A) \times B_3) - A$ keeps only the *missing* pairs: (s2, p4), (s3, p1), (s3, p4), (s4, p1).
- Projecting on **sno** gives the disqualified suppliers: {s2, s3, s4}.
- $\pi_{sno}(A) - \{s2, s3, s4\} = \{s1\}$.

Chapter 3

SQL-1 - SQL Introduction

3.1 Relational terminology

Although we already defined the relational model, let's align our vocabulary with SQL:

- **Relation** translates to a **table**.
- **Tuple** translates to a **row**.
- **Attribute** translates to a **column**.

Throughout this chapter, we use the following example **Students** table:

Students

<u>sid</u>	name	login	age	gpa
53666	Jones	jones@cs	18	3.4
53688	Smith	smith@ee	18	3.2
53777	White	white@cs	19	4.0

Primary Key (PK). The column (or group of columns) that uniquely identifies a row. Two rows **cannot** share the same PK. Above, sid is the primary key.

3.2 Basic SQL queries

3.2.1 The simplest query

To retrieve all contents of a table, use `SELECT *`:

```
SELECT * FROM Students S
```

`S` is an optional *alias* (range variable) for the **Students** table, useful when referencing columns. This returns the full table shown above.

3.2.2 Showing specific columns (Projection)

To retrieve only certain attributes, replace `*` with column names:

```
SELECT S.name, S.login FROM Students S
```

Projection. Selecting specific columns from a table. The query above projects `name` and `login` from `Students`.

name	login
Jones	jones@cs
Smith	smith@ee
White	white@cs

3.2.3 Showing specific rows (Selection)

To filter rows, use the `WHERE` clause:

```
SELECT * FROM Students S WHERE S.age = 18
```

Selection. Filtering rows based on a condition. The query above selects students whose age equals 18.

sid	name	login	age	gpa
53666	Jones	jones@cs	18	3.4
53688	Smith	smith@ee	18	3.2

3.2.4 SQL vocabulary

The general form of a basic SQL query:

```
SELECT [DISTINCT] target-list FROM relation-list WHERE qualification
```

- **relation-list:** a list of table names, each possibly followed by an alias.
- **qualification:** comparisons combined with `AND`, `OR`, `NOT`. Each comparison: `Attr op Const` or `Attr1 op Attr2`.
- **target-list:** a list of attributes from the tables in *relation-list*.
- **DISTINCT:** optional keyword; eliminates duplicate rows. By default, SQL keeps duplicates (the result is a *multiset*).

3.2.5 Evaluation order

Conceptually, a SQL query is evaluated as follows:

1. **FROM:** compute the cross-product of all tables.
2. **WHERE:** discard tuples that fail the condition (*selection*).
3. **SELECT:** keep only the requested columns (*projection*).
4. If **DISTINCT** is specified, eliminate duplicate rows.

3.3 Querying multiple relations

To join two tables, list them both in `FROM` and express the join condition in `WHERE`. Consider the following `Enrolled` table alongside `Students`:

Enrolled

<u>sid</u>	<u>cid</u>	<u>grade</u>
53831	Carnatic101	C
53831	Reggae203	B
53650	Topology112	A
53666	History105	B

“Find all students with grade B”:

```
SELECT S.name, E.cid
FROM Students S, Enrolled E
WHERE S.sid = E.sid AND E.grade = 'B'
```

S.name	E.cid
Jones	History105

3.3.1 Naïve evaluation in steps

The evaluation of the query above proceeds conceptually as follows:

Step 1 – Cross Product (FROM): Combine every row of **Students** with every row of **Enrolled**:

S.sid	S.name	S.login	S.age	S.gpa	E.sid	E.cid	E.grade
53666	Jones	jones@cs	18	3.4	53831	Carnatic101	C
53666	Jones	jones@cs	18	3.4	53831	Reggae203	B
53666	Jones	jones@cs	18	3.4	53650	Topology112	A
53666	Jones	jones@cs	18	3.4	53666	History105	B
53688	Smith	smith@ee	18	3.2	53831	Carnatic101	C
53688	Smith	smith@ee	18	3.2	53831	Reggae203	B
53688	Smith	smith@ee	18	3.2	53650	Topology112	A
53688	Smith	smith@ee	18	3.2	53666	History105	B

Step 2 – Apply the predicate (WHERE): Keep only rows where **S.sid = E.sid and E.grade = 'B'**. Only one row survives: Jones (sid 53666) enrolled in History105 with grade B.

Step 3 – Project the requested columns (SELECT): Keep only **S.name** and **E.cid**, yielding the final result: (Jones, History105).

Running example: Sailors, Boats, Reserves

The following three relations are used throughout the rest of this chapter.

Sailors

<u>sid</u>	sname	rating	age
22	Dustin	7	45.0
31	Lubber	8	55.5
95	Bob	3	63.5

Boats

<u>bid</u>	bname	color
101	Interlake	blue
102	Interlake	red
103	Clipper	green
104	Marine	red

Reserves

<u>sid</u>	<u>bid</u>	day
22	101	10/10/96
95	103	11/12/96

3.4 Aggregate operators

Aggregate function. A function that operates on a set (or multiset) of values from a column and returns a **single** summary value. Aggregates are a significant extension of relational algebra.

Function	Description
COUNT(*)	number of rows
COUNT([DISTINCT] A)	number of (distinct) non-null values in A
SUM([DISTINCT] A)	sum of (distinct) values
AVG([DISTINCT] A)	average of (distinct) values
MAX(A)	maximum value
MIN(A)	minimum value

3.4.1 Examples

Find the number of sailors:

```
SELECT COUNT(*) FROM Sailors S
```

Find the average age of sailors with rating 10:

```
SELECT AVG(S.age) FROM Sailors S WHERE S.rating = 10
```

Count distinct ratings of sailors named Bob:

```
SELECT COUNT(DISTINCT S.rating) FROM Sailors S WHERE S.sname = 'Bob'
```

3.4.2 Name and age of the oldest sailor(s)

This example illustrates a common pitfall with aggregates.

Incorrect query

```
SELECT S.sname, MAX(S.age) FROM Sailors S
```

This is **wrong**: MAX returns a single value over the entire column, so there is no meaningful way to pair it with a particular sname.

Correct approach — subquery

```
SELECT S.sname, S.age
FROM Sailors S
WHERE S.age = (SELECT MAX(S2.age) FROM Sailors S2)
```

The inner query computes the maximum age; the outer query retrieves every sailor matching that value.

Scalar subquery. A subquery that returns exactly one row and one column. It can appear wherever a single value is expected (e.g. in a WHERE comparison).

An equivalent alternative places the subquery on the left:

```
SELECT S.sname, S.age
FROM Sailors S
WHERE (SELECT MAX(S2.age) FROM Sailors S2) = S.age
```

3.5 GROUP BY

So far, aggregates have been applied to *all* qualifying tuples at once. Sometimes we need to apply them to each of several *groups* independently.

Consider: find the age of the youngest sailor for each rating level.

- We don't know in advance how many rating levels exist or what their values are.
- If ratings range from 1 to 10, we *could* write 10 separate queries:

For $i = 1, 2, \dots, 10$:

```
SELECT MIN(S.age) FROM Sailors S WHERE S.rating = i
```

This is obviously impractical.

GROUP BY clause. Partitions tuples into groups based on equal values of the specified columns, then applies aggregate functions independently to each group.

3.5.1 Using GROUP BY

We use an extended `Sailors` table for the following examples:

Sailors (extended)

sid	sname	rating	age
22	Dustin	7	45.0
31	Lubber	8	55.5
95	Bob	3	63.5
52	Smith	3	74.5
68	John	7	56.5
81	Ana	7	71.5

“Find the youngest age for each rating level”:

```
SELECT S.rating, MIN(S.age) FROM Sailors S GROUP BY S.rating
```

The tuples are first partitioned by `rating` (groups: rating 3, rating 7, rating 8), then `MIN(S.age)` is evaluated within each group:

rating	MIN(age)
3	63.5
7	45.0
8	55.5

GROUP BY semantics. The `SELECT` clause may only contain:

- attributes listed in the `GROUP BY` clause, and
- aggregate functions applied to other attributes.

Any non-aggregated column in `SELECT` that is *not* in `GROUP BY` is an error.

3.5.2 HAVING — filtering groups

`HAVING` filters *groups* (after grouping), just as `WHERE` filters *individual tuples* (before grouping).

“Find the youngest age for each rating level that has **at least two** sailors”:

```
SELECT S.rating, MIN(S.age)
FROM Sailors S
GROUP BY S.rating
HAVING COUNT(*) > 1
```

Groups with only one member (rating 8, containing only Lubber) are eliminated:

rating	MIN(age)
3	63.5
7	45.0

HAVING clause. A filter applied *after* grouping. It uses aggregate conditions to keep or discard entire groups, unlike `WHERE` which operates on individual rows before any grouping.